

Assignment_22.1

Task1

Create a kafka topic named KeyLessTopic.

Inside KeyLessTopic insert following data:

```
{"name":"John", "exp":16}
```

```
{"name":"Finn", "exp":20}
```

```
{"name":"Cylin", "exp":18}
```

##Understanding Kafka, starting a broker/Leader that pushes messages to be consumed at consumer.

Demonstrating ways to create multinode replicas with multiple partitions.

Different ways to create a topic(keyed/keyless) based on the key property.

##Formatting messages with custom values of key value separator &

##reading the same consumed messages at consumer using custom format using key separator property.

```
##Command below starts the Zookeeper
```

```
$KAFKA_HOME/bin/zookeeper-server-start.sh $KAFKA_HOME/config/zookeeper.properties
```

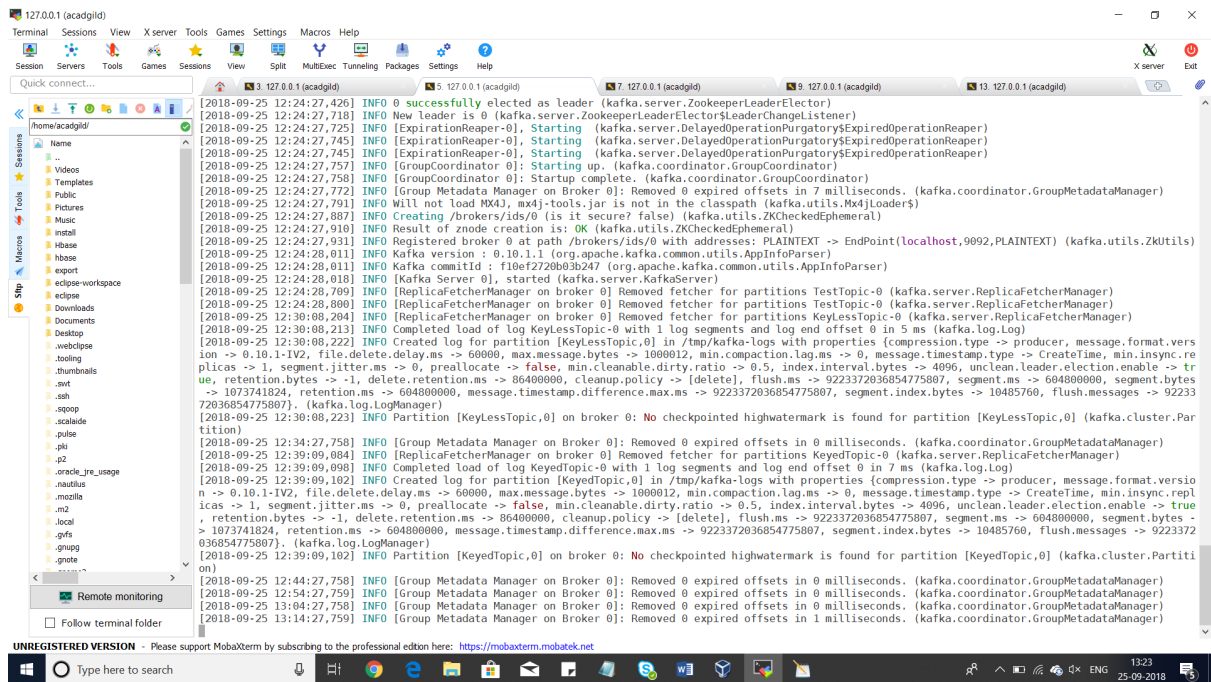
The image shows a terminal window with a title bar that reads "127.0.0.1 (acagild)". The terminal content shows the execution of a script to start ZooKeeper. The output includes a warning about deprecated scripts and a successful startup of the ZooKeeper server. The logs show the server's configuration, including the data directory, log directory, and the number of servers in the quorum. The terminal also shows the startup of the Hadoop and Yarn daemons. A sidebar on the left shows a file explorer view of the /home/acagild/ directory. The top bar displays various application icons and system status.

```
127.0.0.1 (acadgild)
Terminal Sessions View X server Tools Games Settings Macros Help
Quick connect...
home/acadgild/
Name
Videos
Templates
Public
Pictures
Music
install
hbase
export
eclipse-workspace
eclipse
Downloads
Documents
Desktop
webclipse
tooling
thumbnails
.svn
.ssh
.sqoop
.scalade
.pulsee
.gtk
.p2
.oracle_jre_usage
.nautius
.mozila
.m2
.local
.gifs
.grugp
.grote
Remote monitoring
Follow terminal folder
UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: https://mobaxterm.mobatek.net
Type here to search
```

The below command is used to start broker

`$KAFKA_HOME/bin/kafka-server-start.sh $KAFKA_HOME/config/server.properties`

```
127.0.0.1 (acadgild)
Terminal Sessions View X server Tools Games Settings Macros Help
Quick connect...
home/acadgild/
Name
Videos
Templates
Public
Pictures
Music
install
hbase
export
eclipse-workspace
eclipse
Downloads
Documents
Desktop
webclipse
tooling
thumbnails
.svn
.ssh
.sqoop
.scalade
.pulsee
.gtk
.p2
.oracle_jre_usage
.nautius
.mozila
.m2
.local
.gifs
.grugp
.grote
Remote monitoring
Follow terminal folder
UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: https://mobaxterm.mobatek.net
Type here to search
```



##Create a topic named KeyLessTopic wherein partition is set to 1 & replication factor is set to 1. Since this

example is being run over a single node cluster

\$KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyLessTopic --zookeeper localhost:2181 --partitions 1 --replication-factor 1

```
[acagdild@localhost ~]$ $KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyLessTopic --zookeeper localhost:2181 --partitions 1 --replication-factor 1
Created topic "KeyLessTopic".
You have new mail in /var/spool/mail/acagdild
```

Running a console producer. The command initiates a communication at port 9092 with the

topic name KeyLessTopic

\$KAFKA_HOME/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic KeyLessTopic

These are the messages passed as a producer to topic named KeyLessTopic.

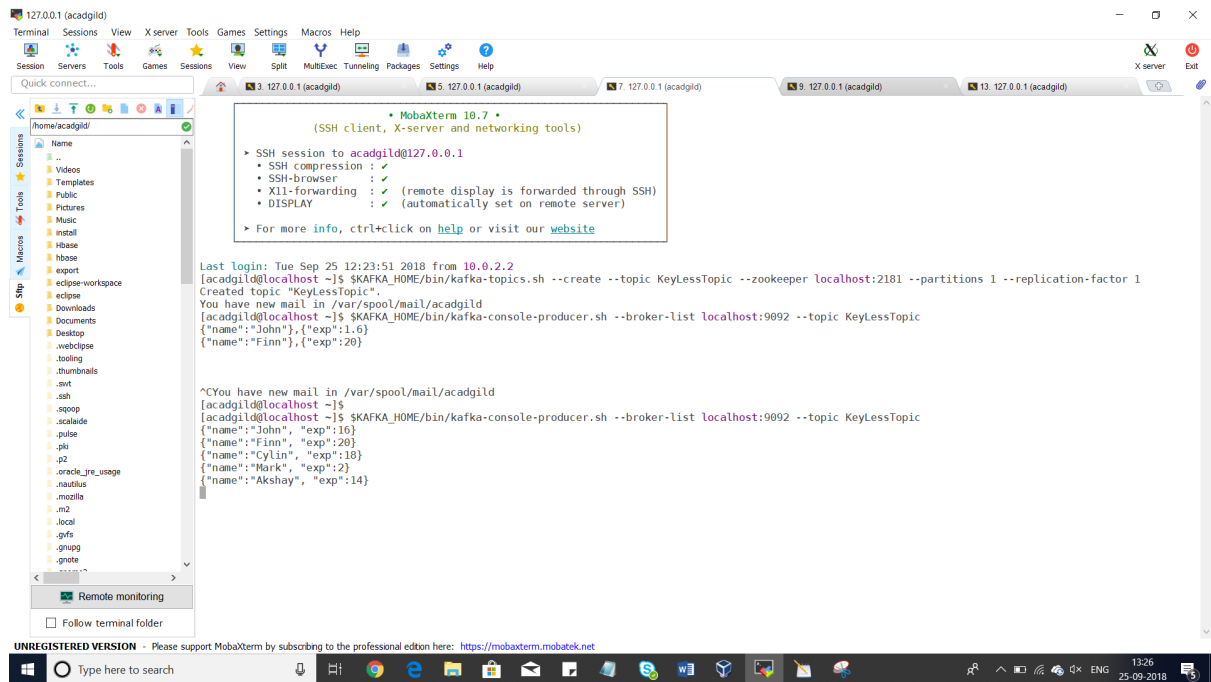
{"name":"John", "exp":16}

{"name":"Finn", "exp":20}

{"name":"Cylin", "exp":18}

{"name":"Mark", "exp":2}

{"name":"Akshay", "exp":14}



Task 2: Create a console consumer that reads KeyLessTopic from beginning

Running a console consumer to read the type KeyLessTopic sent by the producer.

--from-beginning allows the topic messages to be read from beginning. The key property is set to false since no key value is involved.

`$KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyLessTopic --from-beginning --zookeeper localhost:2181 --property print.key=false`

```
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyLessTopic --from-beginning --zookeeper localhost:2181 --property print.key=false
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap
server] instead of [zookeeper].
{"name":"John"}, {"exp":1.6}
{"name":"Finn"}, {"exp":20}
```

```
{"name":"John", "exp":16}
{"name":"Finn", "exp":20}
{"name":"Cylin", "exp":18}
{"name":"Mark", "exp":2}
{"name":"Akshay", "exp":14}
{"name":"John", "exp":16}
{"name":"Finn", "exp":20}
{"name":"Cylin", "exp":18}
{"name":"Mark", "exp":2}
{"name":"Akshay", "exp":14}
```

```
{"name":"John", "exp":16}
{"name":"Finn", "exp":20}
{"name":"Cylin", "exp":18}
{"name":"Mark", "exp":2}
{"name":"Akshay", "exp":14}
```

Task 3:

Create a kafka topic named KeyedTopic. Inside KeyedTopic insert following data:

The part before comma(,) should be treated as key and after comma(,) should be treated as value

```
{"name":"John"},{"exp":16}
```

```
{"name":"Finn"},{"exp":20}
```

```
{"name":"Cylin"},{"exp":18}
```

```
{"name":"Mark"},{"exp":2}
```

```
{"name":"Akshay"},{"exp":14}
```

##Create a topic named KeyedTopic

```
$KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyedTopic --zookeeper
```

```
localhost:2181 --partitions 1 --replication-factor 1
```

```
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-topics.sh --create --topic KeyedTopic --zookeeper localhost:2181 --partitions 1 --replication-factor 1
Created topic "KeyedTopic".
You have new mail in /var/spool/mail/acadgild
```

Running a console producer

```
$KAFKA_HOME/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic
KeyedTopic --property parse.key=true --property key.separator=","
```

```
{"name":"John"},{"exp":16}
```

```
{"name":"Finn"},{"exp":20}
```

```
{"name":"Cylin"},{"exp":18}
```

```
{"name":"Mark"},{"exp":2}
```

```
{"name":"Akshay"},{"exp":14}
```

```
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic KeyedTopic --property parse.key=true --property key.separator=","
{"name":"John"},{"exp":16}
{"name":"Finn"},{"exp":20}
{"name":"Cylin"},{"exp":18}
{"name":"Mark"},{"exp":2}
{"name":"Akshay"},{"exp":14}
```

Task 4:

Create a console consumer that reads KeyedTopic from beginning

The key and value should be separated by '-'

Ans: ## Running a console consumer to read the type KeyedTopic sent by the producer.
\$KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyedTopic --frombeginning
--zookeeper localhost:2181 --property print.key=true --property key.separator="-"

```
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic KeyedTopic --from-beginning --zookeeper localhost:2181 --property print.key=true --pr
operty key.separator="-"
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-
server] instead of [zookeeper].
{"name":"John"}-{"exp":16}
{"name":"Finn"}-{"exp":20}
{"name":"Cylind"}-{"exp":18}
{"name":"Mark"}-{"exp":2}
{"name":"Akshay"}-{"exp":14}
```

Note: From the screen shot we can see that key.separator="," got replaced by
key.separator="-"